

Concurrent Programming Problems

OS

Spring 2011



Concurrency pros and cons

- Concurrency is good for users
 - One of the reasons for multiprogramming
 - Working on the same problem, simultaneous execution of programs, background execution
- Concurrency is a “pain in the neck” for the system
 - Access to shared data structures
 - Danger of deadlock due to resource contention

Linked list example

`insert_after(current,new): new->next=current.next`

`current.next=new`





Linked list example




remove_next(current):

```
tmp=current.next;
```

```
current.next=current.next.next;
```

```
free(tmp);
```





Mutual Exclusion

- OS is an instance of concurrent programming
 - Multiple activities may take place at the ‘same time’
- Concurrent execution of operations involving multiple steps is problematic
 - Example: updating linked list
- Concurrent access to a shared data structure must be mutually exclusive

Atomic Operations

- A generic solution is to ensure *atomic* execution of operations
 - All the steps are *perceived* as executed in a single point of time

insert_after(current,new) remove_next(current), or

remove_next(current) insert_after(current,new)

The Critical Section Problem

- n processes P_0, \dots, P_{n-1}
- No assumptions on relative process speeds, no synchronized clocks, etc...
 - Models inherent non-determinism of process scheduling
- No assumptions on process activity when executing within critical section and remainder sections
- **The problem:** Implement a general mechanism for entering and leaving a critical section.


Success Criteria

1. *Mutual exclusion*: Only one process is in the critical section at a time.
2. *Progress*: There is no deadlock: some process will eventually get into the critical section.
3. *Fairness*: There is no starvation: no process will wait indefinitely while other processes continuously enter the critical section.
4. *Generality*: It works for N processes.

Peterson's Algorithm

- There are two shared arrays, initialized to 0.
 - **q[n]**: q[i] is the stage of process i in entering the CS. Stage n means the process is in the CS.
 - **turn[n-1]**: turn[j] says which process has to wait in stage j.
- The algorithm for process i:

```
for (j=1; j<n; j++) {  
    q[i] = j;  
    turn[j] = i;  
    while (( $\exists k \neq i$  s.t.  $q[k] \geq j$ )  
           && (turn[j] == i)) {  
        skip;  
    }  
}  
critical section  
q[i] = 0;
```



Proof of Peterson's algorithm

```
for (j=1; j<n; j++) {  
    q[i] = j;  
    turn[j] = i;  
    while(( $\exists k \neq i$  st  $q[k] \geq j$ )  
        && (turn[j] == i)) {  
        skip;  
    }  
}  
critical section  
q[i] = 0;
```

Definition: Process a is **ahead of** process b if $q[a] > q[b]$.

Lemma 1:

A process that is ahead of all others advances to the next stage (increments $q[i]$).

Proof: This process is not stuck in the while loop because the first condition does not hold.

Proof of Peterson's algorithm

```
for (j=1; j<n; j++) {  
    q[i] = j;  
    turn[j] = i;  
    while(( $\exists k \neq i$  st  $q[k] \geq j$ )  
        && (turn[j] == i)) {  
        skip;  
    }  
}  
critical section  
q[i] = 0;
```

Lemma 2:

When a process advances to stage $j+1$, if it is not ahead of all processes, then there is at least one other process at stage j .

Proof:

To exit the while loop another process had to take turn[j].

Proof of Peterson's algorithm

```
for (j=1; j<n; j++) {  
    q[i] = j;  
    turn[j] = i;  
    while(( $\exists k \neq i$  st  $q[k] \geq j$ )  
        && (turn[j] == i)) {  
        skip;  
    }  
}  
critical section  
q[i] = 0;
```

Lemma 3:

If there is more than one process at stage j , then there is at least one process at every stage below j .

Proof:

Use lemma 2, and prove by induction on the stages.

Proof of Peterson's algorithm

```
for (j=1; j<n; j++) {  
    q[i] = j;  
    turn[j] = i;  
    while(( $\exists k \neq i$  st  $q[k] \geq j$ )  
        && (turn[j] == i)){  
        skip;  
    }  
}  
critical section  
q[i] = 0;
```

Lemma 4:

The maximal number of processes at stage j is $n-j+1$

Proof:

By lemma 3, there are at least $j-1$ processes at lower stages.

Proof of Peterson's algorithm

- Mutual Exclusion:
 - By Lemma 4, only one process can be in stage n
- Progress:
 - There is always at least one process that can advance:
 - If a process is ahead of all others it can advance
 - If no process is ahead of all others, then there is more than one process at the top stage, and one of them can advance.
- Fairness:
 - A process will be passed over no more than n times by each of the other processes (proof: in the notes).
- Generality:
 - The algorithm works for any n .

Peterson's Algorithm

- Peterson's algorithm creates a critical section mechanism without any help from the OS.
- All the success criteria hold for this algorithm.
- It does use busy wait (no other option).

Classical Problems of Synchronization

Classical Problems of Synchronization

- Bounded-Buffer Problem
- Readers and Writers Problem
- Dining-Philosophers Problem

Bounded-Buffer Problem

- One cyclic buffer that can hold up to N items
- **Producer** and **consumer** use the buffer
 - The buffer absorbs fluctuations in rates.
- The buffer is shared, so protection is required.
- We use **counting semaphores**:
 - the number in the semaphore represents the number of resources of some type

Bounded-Buffer Problem

- Semaphore **mutex** initialized to the value 1
 - Protects the index into the buffer
- Semaphore **full** initialized to the value 0
 - Indicates how many items in the buffer are full (can read them)
- Semaphore **empty** initialized to the value N
 - Indicates how many items in the buffer are empty (can write into them)

Bounded-Buffer Problem – Cont.

Producer:

```
while (true) {  
    produce an item  
    P (empty);  
    P (mutex);  
    add the item to the buffer  
    V (mutex);  
    V (full);  
}
```

Consumer:

```
while (true) {  
    P (full);  
    P (mutex);  
    remove an item from buffer  
    V (mutex);  
    V (empty);  
    consume the item  
}
```

Readers-Writers Problem

- A data structure is shared among a number of concurrent processes:
 - Readers – Only read the data; They do not perform updates.
 - Writers – Can both read and write.
- Problem
 - Allow multiple readers to read at the same time.
 - Only one writer can access the shared data at the same time.
 - If a writer is writing to the data structure, no reader may read it.

Readers-Writers Problem:

First Solution

- Shared Data:
 - The data structure
 - Integer **readcount** initialized to 0.
 - Semaphore **mutex** initialized to 1.
 - Protects readcount
 - Semaphore **write** initialized to 1.
 - Makes sure the writer doesn't use data at the same time as any readers

Readers-Writers Problem:

First solution

- The structure of a **writer** process:

```
while (true) {  
    P (write) ;  
  
    writing is performed  
  
    V (write) ;  
}
```


Readers-Writers Problem:

First solution

- The structure of a **reader** process:

```
while (true) {  
    P (mutex) ;  
    readcount ++ ;  
    if (readcount == 1)  
        P (write) ;  
    V (mutex)  
    reading is performed  
    P (mutex) ;  
    readcount - - ;  
    if (readcount == 0)  
        V (write) ;  
    V (mutex) ;  
}
```

Readers-Writers Problem:

First solution

- The structure of a **reader** process:

```
while (true) {  
    P (mutex) ;  
    readcount ++ ;  
    if (readcount == 1)  
        P (write) ;  
    V (mutex)  
    reading is performed  
    P (mutex) ;  
    readcount - - ;  
    if (readcount == 0)  
        V (write) ;  
    V (mutex) ;  
}
```

This solution is
not perfect:

What if a writer is
waiting to write
but there are
readers that read
all the time?

Writers are subject to
starvation!

Second solution: Writer Priority

- Extra semaphores and variables:
 - Semaphore **read** initialized to 1 – inhibits readers when a writer wants to write.
 - Integer **writcount** initialized to 0 – counts waiting writers.
 - Semaphore **write_mutex** initialized to 1 – controls the updating of writcount.
 - Semaphore mutex now called **read_mutex**
 - **Queue** semaphore used only in the reader

Second solution: Writer Priority

The writer:

```
while (true) {  
    P(write_mutex)  
    writecount++; //counts number of waiting writers  
    if (write_count == 1)  
        P(read)  
    V(write_mutex)  
  
    P (write) ;  
    writing is performed  
    V(write) ;  
  
    P(write_mutex)  
    writecount--;  
    if (writecount == 0)  
        V(read)  
    V(write_mutex)  
}
```


Second Solution: Writer Priority (cont.)

The reader:

```
while (true) {  
    P(queue)  
    P(read)  
    P(read_mutex) ;  
    readcount ++ ;  
    if (readcount == 1)  
        P(write) ;  
    V(read_mutex)  
    V(read)  
    V(queue)  
    reading is performed  
    P(read_mutex) ;  
    readcount - - ;  
    if (readcount == 0)  
        V(write) ;  
    V(read_mutex) ;  
}
```

Queue semaphore, initialized to 1:
Since we don't want to allow more than one reader at a time in this section (otherwise the writer will be blocked by multiple readers when doing P(read).)

Dining-Philosophers Problem



Shared data

Bowl of rice (data set)

Semaphore **chopstick** [5] initialized to 1

Dining-Philosophers Problem – Cont.

- The structure of Philosopher i:

```
While (true) {  
    P ( chopstick[i] );  
    P ( chopstick[ (i + 1) % 5] );  
  
    eat  
  
    V ( chopstick[i] );  
    V ( chopstick[ (i + 1) % 5] );  
  
    think  
}
```

- This can cause deadlocks ☹️

Dining Philosophers Problem

- This abstract problem demonstrates some fundamental limitations of deadlock-free synchronization.
- There is no symmetric solution
 - Any symmetric algorithm leads to a symmetric output, that is everyone eats (which is impossible) or no-one does.

Possible Solutions

- Use a waiter
- Execute different code for odd/even
- Give them another chopstick
- Allow at most 4 philosophers at the table
- Randomized (Lehmann-Rabin)

Summary

- Concurrency can cause serious problems if not handled correctly.
 - Atomic operations
 - Critical sections
 - Semaphores and mutexes
 - Careful design to avoid deadlocks and livelocks.